

Qwen 3.8 27B available on Cerebras at 1500 tokens/s

From: Hacker News

Source: Hacker News

The Brief

****What shipped****

Qwen 3.8 27B model is available on Cerebras.

****Capabilities****

The model has a processing speed of 1500 tokens per second. Specifics on modalities, context length, or benchmark results are not stated.

****Availability****

The model can be accessed on the Cerebras platform. Pricing and tier details are not stated.

****Why it matters****

This release may offer enhanced processing speed for applications using the Qwen model, potentially benefiting developers looking for efficient AI solutions.

****Caveats****

Details on performance benchmarks, such as real-world effectiveness or comparison to existing models, are not provided. Claims regarding speed have not been independently verified.

Full Announcement Details

Qwen 3.8 27B model is available on Cerebras. The model has a processing speed of 1500 tokens per second. Specifics on modalities, context length, or benchmark results are not stated.

Want to learn more?

<https://inference-docs.cerebras.ai/models/overview>

Daily AI Dev Digest

Thursday, September 03, 2026

This digest tracks the latest developments from the frontier AI labs - model releases, product and API launches, research reports, and capability results - as they ship.